

## 1. Autoencoders

What does a deep neural network do?


It learns important feature from the input.

Important features that allow to do specific task on the data (eg. classification, regression, generalization).

Autoencoders are Neural networks that works in an unsupervised manner.

\* no need labeled data.

## 2. Autoencoder detail



↑ maximize the similarity between input & output.

↓ minimize the reconstruction error.

## 3. GAE theory.

one convolutional Graph neural network:


- produces a low dimensional embedding representation

$$\bar{X} = \text{GCN}(A, X) = \text{ReLU}(\tilde{A}XW_0)$$

$$\text{with } \tilde{A} = D^{-\frac{1}{2}} A D^{-\frac{1}{2}},$$

} Encoder.

$Z = \bar{X} \Rightarrow$  Embedding latent space.



use inner product to construct the adjacency matrix  $\hat{A}$ .

Inner product between latent variable  $z$ .


$$\hat{A} = \text{logistic sigmoid}(zz^T)$$

## Jupyter Notebook:

①

```
class GCNEncoder(torch.nn.Module):  
    def __init__(self, in_channels, out_channels):  
        super(GCNEncoder, self).__init__()  $\Rightarrow$  initialize  
        self.conv1 = GCNConv(in_channels, 2*out_channels, cached = True)  
        self.conv2 = GCNConv(2*out_channels, out_channels, cached = True)  
  
    def forward(self, x, edge_index):  
        x = self.conv1(x, edge_index).relu()  
        return self.conv2(x, edge_index)
```

$\Rightarrow$  transductive learning



② train & test function.

```
def train():  
    model.train()  
    optimizer.zero_grad()  $\Rightarrow$  clears old gradients from the last step OR they will accumulate.  
    z = model.encode(x, train_pos_edge_index)  $\Rightarrow$  embedding the input data  
    loss = model.recon_loss(z, train_pos_edge_index)  $\Rightarrow$  calculate the reconstruct loss  
    loss.backward()  $\Rightarrow$  doing backpropagation to compute the gradient loss with respect to all model parameters  
    optimizer.step()  $\Rightarrow$  update the model parameter.  
    return float(loss)  
  
def test(pos_edge_index, neg_edge_index):  
    model.eval()  
    with torch.no_grad():  
        z = model.encode(x, train_pos_edge_index)  
    return model.test(z, pos_edge_index, neg_edge_index)
```

③ . test the result.

```
for epoch in range(1, epochs+1):  
    loss = train()  
  
    auc, ap = test(data.test_pos_edge_index, data.test_neg_edge_index)  
    print('Epoch: {:03d}, AUC: {:.4f}, AP: {:.4f}'.format(epoch, auc, ap))
```

④ Use tensorboard to compare the results rather than AUC & AP.



store the results :

```
writer = SummaryWriter('runs/GAE_experiment' + '2d_100_epochs')  
  
for epoch in range(1, epochs+1):  
    loss = train()  
    auc, ap = test(data.test_pos_edge_index, data.test_neg_edge_index)  
    print('Epoch: {:03d}, AUC: {:.4f}, AP: {:.4f}'.format(epoch, auc, ap))  
  
    writer.add_scalar('auc train', auc, epoch)  
    writer.add_scalar('ap train', ap, epoch)
```

Then we can compare different results in Tensorboard.

different dim, different epochs.

## 4. Variational Autoencoders



So we need to let  $p(z|x)$  and  $q(z|x)$

**KL-Divergence**  $\Rightarrow$  measures the distance between distribution.

$$KL(q(z|x) || p(z|x))$$

$$\text{Loss} = -\mathbb{E}_{p(z|x)} \log q(z|x) + KL(p(z|x) || q(z))$$

**Variational Lower Bound:**

How well the network is able to reconstruct the input?


(GVAE)

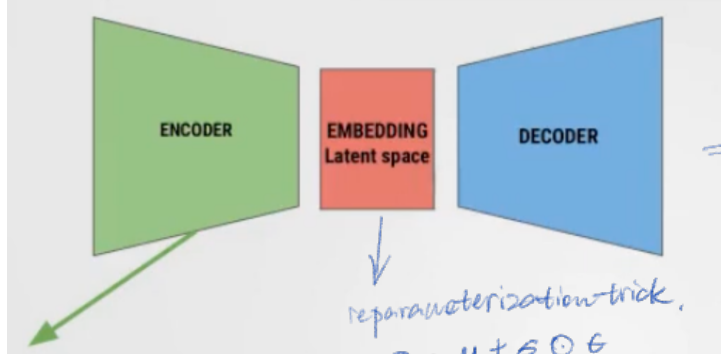
★ Reparameterization trick.

$$z = \mu + \sigma \odot \epsilon \quad \epsilon \sim \text{Norm}(0,1)$$

Original form

Reparameterised form





⇒ Still using inner product.

$$\hat{A} = \text{logistic sigmoid}(ZZ^T)$$

Two convolutional GNN.

$$z = u + \sigma \odot \epsilon$$

$$\epsilon \sim \text{Norm}(0,1)$$

GCN1: produces an lower dimensioned embedding

$$\bar{x} = \text{GCN}(A, x) = \text{ReLU}(\hat{A}xW_0) \text{ with } \hat{A} = \bar{D}^{-\frac{1}{2}}A\bar{D}^{-\frac{1}{2}}$$

GCN2: generates  $\mu$  and  $\log \sigma^2$ .

$$\mu = \text{GCN}_\mu(x, A) = \tilde{A}\bar{x}W_1$$

$$\log \sigma^2 = \text{GCN}_\sigma(x, A) = \tilde{A}\bar{x}W_1$$

Jupyter Notebook.

```
dataset = Planetoid("../", "CiteSeer", transform=T.NormalizeFeatures())
data = dataset[0]
data.train_mask = data.val_mask = data.test_mask = None
data = train_test_split_edges(data)

class VariationalGCNEncoder(torch.nn.Module):
    def __init__(self, in_channels, out_channels):
        super(VariationalGCNEncoder, self).__init__()
        self.conv1 = GCNConv(in_channels, 2*out_channels, cached = True)
        self.conv_mu = GCNConv(2*out_channels, out_channels, cached = True)
        self.conv_logstd = GCNConv(2*out_channels, out_channels, cached = True)

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index).relu()
        return self.conv_mu(x, edge_index), self.conv_logstd(x, edge_index) #returning two outputs
```

should return  $\mu$  and  $\sigma^2$ .

```
out_channels = 2
num_features = dataset.num_features
epochs = 300

#new line ↗ Apply VGAE model.
model = VGAE(VariationalGCNEncoder(num_features, out_channels))

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)
x = data.x.to(device)
train_pos_edge_index = data.train_pos_edge_index.to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.1)
```

```
def train():
    model.train()
    optimizer.zero_grad()
    z = model.encode(x, train_pos_edge_index)
    loss = model.recon_loss(z, train_pos_edge_index)
    #new line
    loss = loss + (1/data.num_nodes)*model.kl_loss()
    loss.backward()
    optimizer.step()
    return float(loss)

def test(pos_edge_index, neg_edge_index):
    model.eval()
    with torch.no_grad():
        z = model.encode(x, train_pos_edge_index)
    return model.test(z, pos_edge_index, neg_edge_index)
```

Recall new loss function.

$$\text{Loss} = -\mathbb{E}_{p(z|x)} \log q(z|x) + \text{KL}(p(z|x) || q(z))$$

↑  
regularization part.